

Data Analytics III

- Implement Simple Naïve Bayes classification algorithm using Python/R on iris.csv dataset.
- Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset.

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from sklearn import datasets
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import make_scorer, accuracy_score, precision_score, recall_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score, precision_score, recall_score
```

Loading Data set

```
In [3]: # Load the iris dataset
df = pd.read_csv('Iris.csv')
df.head()
```

```
Out [3]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
In [4]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
# 0   Id              150 non-null    int64
# 1   SepalLengthCm   150 non-null    float64
# 2   SepalWidthCm    150 non-null    float64
# 3   PetalLengthCm   150 non-null    float64
# 4   PetalWidthCm    150 non-null    float64
# 5   Species         150 non-null    object
```

```

---  -----
0   Id          150 non-null    int64
1   SepalLengthCm 150 non-null    float64
2   SepalWidthCm  150 non-null    float64
3   PetalLengthCm 150 non-null    float64
4   PetalWidthCm  150 non-null    float64
5   Species       150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB

```

In [5]: `df.isnull().sum`

Out [5]: <bound method NDFrame._add_numeric_operations.<locals>.sum of Id
SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
0 False False False False False
False
1 False False False False False
False
2 False False False False False
False
3 False False False False False
False
4 False False False False False
False
..
...
145 False False False False False
False
146 False False False False False
False
147 False False False False False
False
148 False False False False False
False
149 False False False False False
False
[150 rows x 6 columns]>

In [6]: `df = df.drop(columns= ['Id'])`
`df.head()`

Out [6]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

In [7]: `df.describe()`

Out [7]:


	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

In [8]: `df['Species'].value_counts()`

Out [8]: Iris-setosa 50
Iris-versicolor 50
Iris-virginica 50
Name: Species, dtype: int64


In [9]: `df['SepalLengthCm'].hist()`

Out [9]: <matplotlib.axes._subplots.AxesSubplot at 0x7fdb8e2e2e20>




In [10]: `df['SepalWidthCm'].hist()`

Out [10]: <matplotlib.axes._subplots.AxesSubplot at 0x7fdb8623e670>




```
In [11]: df['PetalLengthCm'].hist()
```

Out [11]: <matplotlib.axes._subplots.AxesSubplot at 0x7fdb85d06340>



```
In [12]: df['PetalWidthCm'].hist()
```


Out [12]: <matplotlib.axes._subplots.AxesSubplot at 0x7fdb85c84e50>



```
In [13]: # Scatterplot
colors = ['red', 'orange', 'blue']
species = ['Iris-virginica', 'Iris-versicolor', 'Iris-setosa']
```

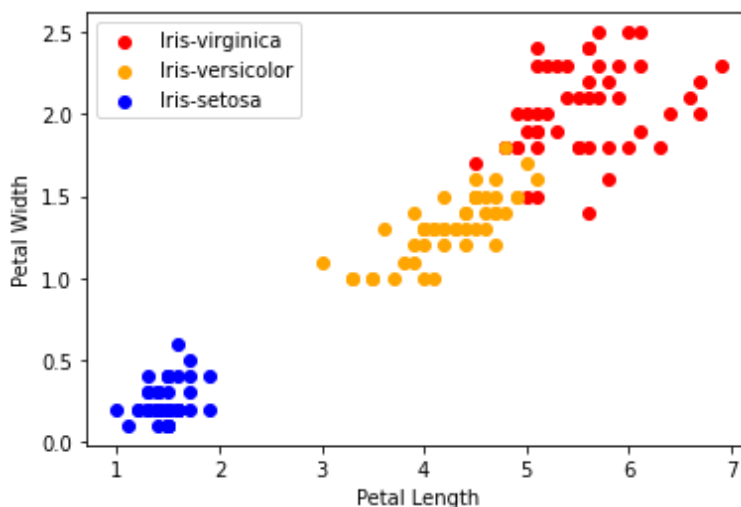
```
In [14]: for i in range(3):
          x = df[df['Species'] == species[i]]
          plt.scatter(x['SepalLengthCm'], x['SepalWidthCm'], c = co
          plt.xlabel("Sepal Length")
          plt.ylabel("Sepal Width")
          plt.legend()
```

Out [14]: <matplotlib.legend.Legend at 0x7fdb85bf7790>



```
In [15]: for i in range(3):
          x = df[df['Species'] == species[i]]
          plt.scatter(x['PetalLengthCm'], x['PetalWidthCm'], c = co
          plt.xlabel("Petal Length")
          plt.ylabel("Petal Width")
          plt.legend()
```

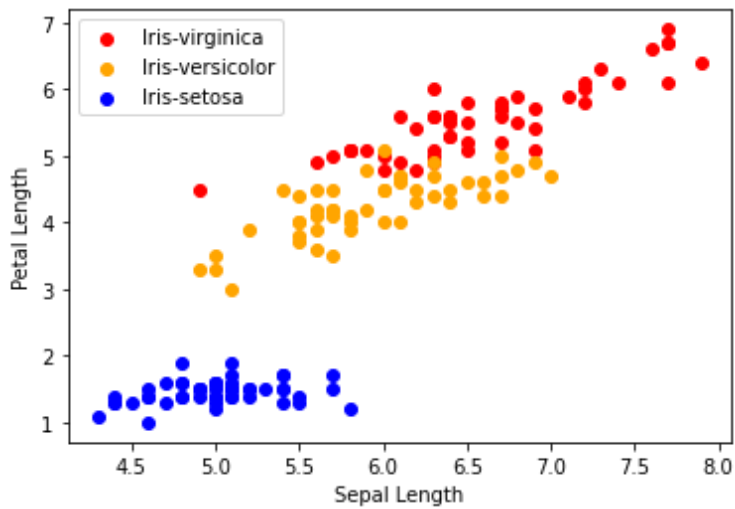
Out [15]: <matplotlib.legend.Legend at 0x7fdb85b78820>



```
In [16]: for i in range(3):
          x = df[df['Species'] == species[i]]
          plt.scatter(x['SepalLengthCm'], x['PetalLengthCm'], c = c
          plt.xlabel("Sepal Length")
```

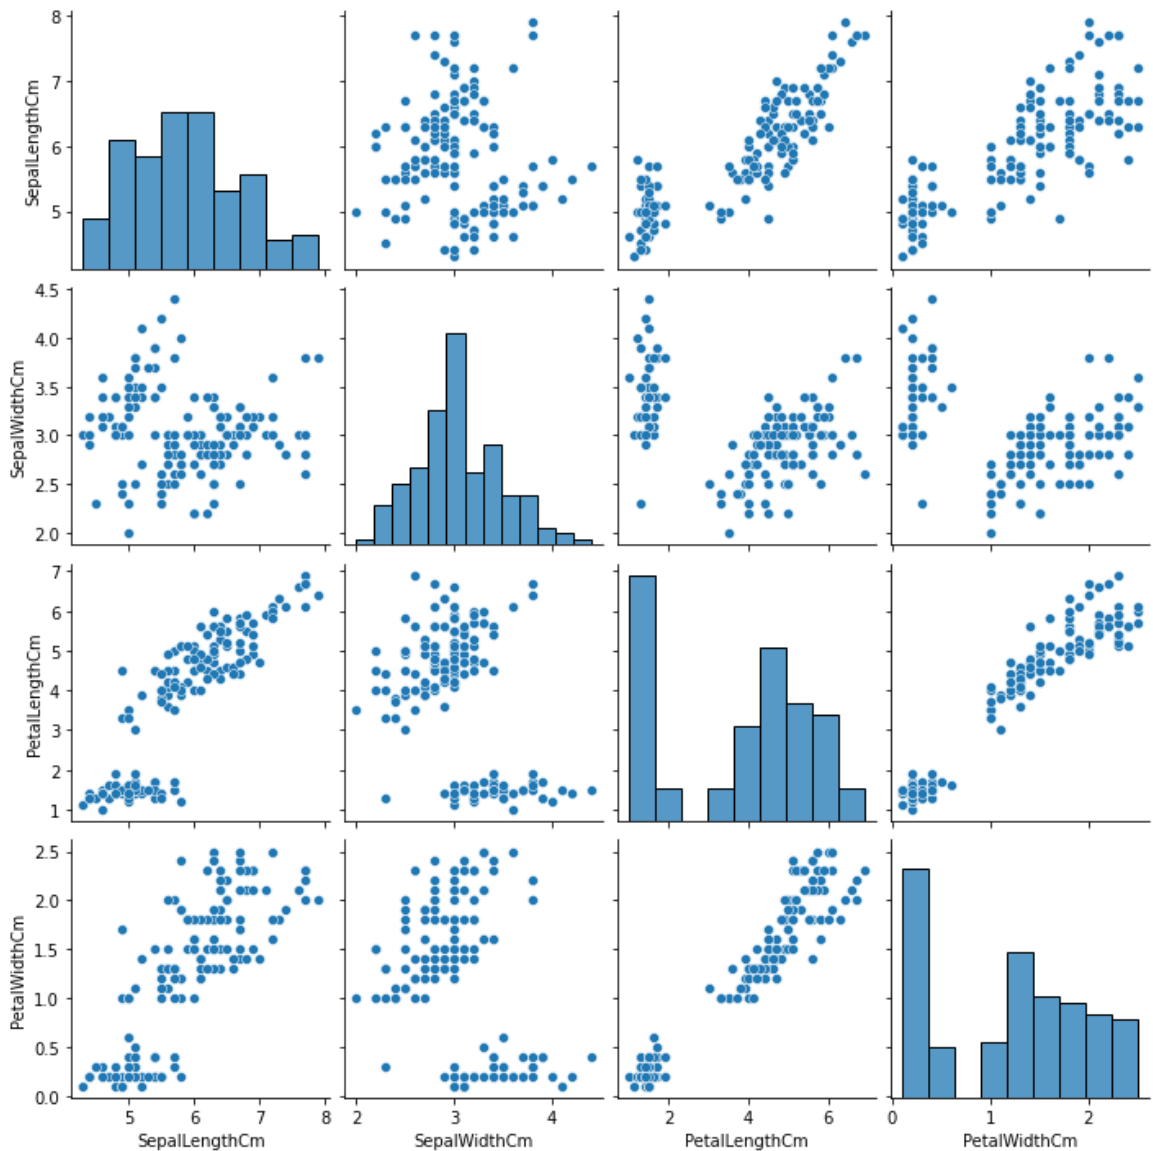
```
plt.ylabel("Petal Length")
plt.legend()
```

Out [16]: <matplotlib.legend.Legend at 0x7fdb85c805b0>



In [17]: `sns.pairplot(df)`

Out [17]: <seaborn.axisgrid.PairGrid at 0x7fdb85c03d30>



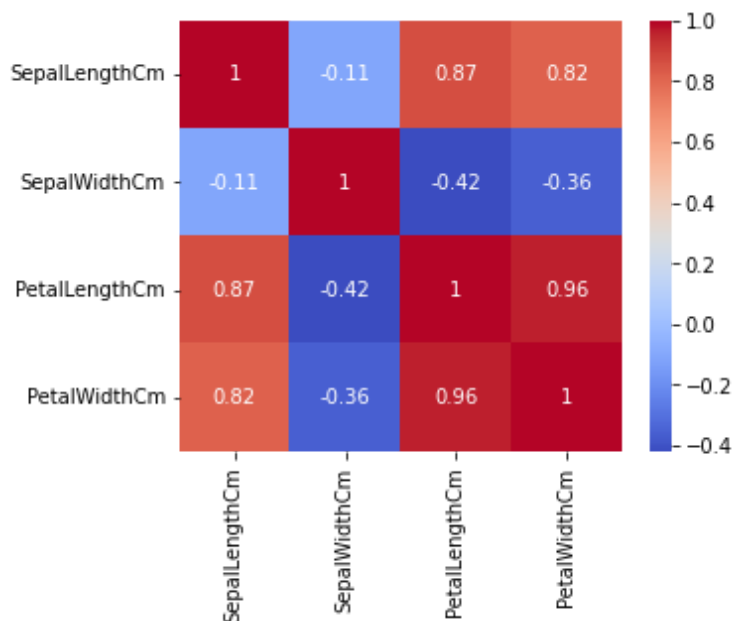
```
In [18]: df.corr()
```

```
Out [18]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
SepalLengthCm	1.000000	-0.109369	0.871754	0.817954
SepalWidthCm	-0.109369	1.000000	-0.420516	-0.356544
PetalLengthCm	0.871754	-0.420516	1.000000	0.962757
PetalWidthCm	0.817954	-0.356544	0.962757	1.000000

```
In [19]: corr = df.corr()  
fig, ax = plt.subplots(figsize=(5,4))  
sns.heatmap(corr, annot=True, ax=ax, cmap = 'coolwarm')
```

```
Out [19]: <matplotlib.axes._subplots.AxesSubplot at 0x7fdb85345820>
```



```
In [20]: from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()
```

```
In [21]: df['Species'] = le.fit_transform(df['Species'])  
df.head()
```

```
Out [21]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
In [33]: X = df.iloc[:, :-1]
```

```
In [34]: X
```

```
Out [34]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
--	---------------	--------------	---------------	--------------

0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2
...
145	6.7	3.0	5.2	2.3
146	6.3	2.5	5.0	1.9
147	6.5	3.0	5.2	2.0
148	6.2	3.4	5.4	2.3
149	5.9	3.0	5.1	1.8

150 rows × 4 columns

```
In [22]: from sklearn.model_selection import train_test_split
# train - 70
# test - 30
X = df.drop(columns=['Species'])
Y = df['Species']
x_train, x_test, y_train, y_test = train_test_split(X, Y, tes
```

```
In [25]: X
```

```
Out [25]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
--	---------------	--------------	---------------	--------------

0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2
...
145	6.7	3.0	5.2	2.3
146	6.3	2.5	5.0	1.9

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
147	6.5	3.0	5.2	2.0
148	6.2	3.4	5.4	2.3
149	5.9	3.0	5.1	1.8

150 rows × 4 columns

In [23]:

```
Y
```

Out [23]:

```
0      0
1      0
2      0
3      0
4      0
..
145    2
146    2
147    2
148    2
149    2
```

Name: Species, Length: 150, dtype: int64

In [35]:

```
gaussian = GaussianNB()
gaussian.fit(x_train, y_train)
Y_pred = gaussian.predict(x_test)
```

In [36]:

```
accuracy_nb=round(accuracy_score(y_test,Y_pred)* 100, 2)
acc_gaussian = round(gaussian.score(x_train, y_train) * 100,

cm = confusion_matrix(y_test, Y_pred)
accuracy = accuracy_score(y_test,Y_pred)
precision =precision_score(y_test, Y_pred,average='micro')
recall = recall_score(y_test, Y_pred,average='micro')
f1 = f1_score(y_test,Y_pred,average='micro')
print('Confusion matrix for Naive Bayes\n',cm)
print('accuracy_Naive Bayes: %.3f' %accuracy)
print('precision_Naive Bayes: %.3f' %precision)
print('recall_Naive Bayes: %.3f' %recall)
print('f1-score_Naive Bayes : %.3f' %f1)
```

Confusion matrix for Naive Bayes

```
[[15  0  0]
 [ 0 11  0]
 [ 0  2 17]]
accuracy_Naive Bayes: 0.956
precision_Naive Bayes: 0.956
recall_Naive Bayes: 0.956
f1-score_Naive Bayes : 0.956
```

In []:

In []:

In []: